

Bayesian inference for infectious disease modelling

Summer School Halle 2022

Fabienne Krauer (LSHTM)



centre for
mathematical
modelling of
infectious
diseases

LONDON
SCHOOL of
HYGIENE
& TROPICAL
MEDICINE



Recommended Readings

Literature

Bayesian Data Analysis (Third edition).

Andrew Gelman, John Carlin, Hal Stern, David Dunson, Aki Vehtari, and Donald Rubin.

Online here: <https://users.aalto.fi/~ave/BDA3.pdf>

Statistical Rethinking (Second edition). Richard McElreath.

Bayesian workflow. ArXiv 2020. Andrew Gelman, Aki Vehtari, Daniel Simpson, Charles C. Margossian, Bob Carpenter, Yuling Yao, Lauren Kennedy, Jonah Gabry, Paul-Christian Bürkner, Martin Modrák. <https://arxiv.org/abs/2011.01808>



Online resources

- https://avehtari.github.io/BDA_course_Aalto/
- <http://florianhartig.github.io/LearningBayes/>
- Stan user guide <https://mc-stan.org/docs/stan-users-guide/index.html>
- MCMC samplers: https://m-clark.github.io/docs/ld_mcmc/
- https://betanalpha.github.io/assets/case_studies/principled_bayesian_workflow.html
- Guided example for model fitting in Stan/rstan: https://mc-stan.org/users/documentation/case-studies/boarding_school_case_study.html
- <https://chi-feng.github.io/mcmc-demo/app.html?algorithm=NaiveNUTS&target=banana>
- https://m-clark.github.io/docs/ld_mcmc/


Table of content

- Inference and Likelihood
- Bayesian theorem and inference
- Sampling algorithms
- The Bayesian workflow
- R packages for fitting ODE models

Linking the model and data: likelihood-based inference



Measures the uncertainty in the outcome of interest from the process model



“likelihood”

“The **likelihood function** (often simply called the likelihood) describes the **joint probability** of the **observed data** as a function of the parameters of the chosen statistical model”

Casella, Statistical Inference, 2002

$$\Pr(Data \mid parameters) = \Pr(y|\theta)$$

The likelihood is NOT a probability density function in the strict sense (i.e. does not integrate to 1)

Likelihood in practice

1) Choose a likelihood function that describes the measurement error best:

Data type	Data example	Likelihood function	Property	PMF/PDF	R function
Count data	Incidence	Poisson	Mean = Variance	$\Pr(X=k) = \frac{\lambda^k e^{-\lambda}}{k!}$	dpois(x, lambda) rpois(n, lambda)
		Negative Binomial	Mean < Variance (=overdispersion)	$\Pr(X = k) = \binom{k+r-1}{r-1} (1-p)^k p^r$	dnbinom(x, size, mu) rnbinom(1, size, mu)
Proportions	Sero-prevalence	Binomial		$\Pr(X = k) = \binom{n}{k} p^k (1-p)^{n-k}$	dbinom(x, size, prob) rbinom(1, size, prob)

Likelihood in practice

2) Apply the chosen likelihood function to a data point and the corresponding model point.

Example: Poisson distribution

At $t = 56$

Observation (“ x ”) = 87 cases

Model prediction (“ λ ”) = 217.9 cases


“by hand”:

$$\Pr(X=k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

$P(\text{data point} \mid \text{theta}) = 2.963404\text{e-}24$

In R:

```
> dpois(87, 217.9)
[1] 2.963404e-24
```



Likelihood in practice


3) Log transform the resulting probability (In R: `dpois(87, 217.9, log=TRUE)`)

4) Repeat this for every data point / model point and **sum** the resulting probabilities (“joint probability” remember?), which yields the **log likelihood**:

$$\log(\Pr(data|\theta)) = \sum_i \log(p(data\ point\ i\ |\theta))$$

$$\Pr(data|\theta) = \prod_i p(data\ point\ i\ |\theta)$$

Log likelihood
-3518.18




Linking the likelihood and inference


What are the log likelihood values if we vary these parameters? → “by hand”: likelihood profile

Example: ODE model with $\beta=0.6$ and $\sigma=0.2$

Fix σ , vary β from 0 to 1





Fix β , vary σ from 0 to 1




Linking the likelihood and inference

Inference explores the log likelihood landscape (surface) while varying all the parameters of interest (the fitted parameters)



Likelihood derived inference methods

	Null hypothesis testing	Maximum likelihood estimation (MLE)	Bayesian inference
Uncertainty	data	data	parameters
output	Rejection of H0 or not	MLE point estimator (set of theta that maximizes the likelihood)	Posterior distribution
Confidence	P-value	Confidence intervals	Credible intervals



Frequentist statistics

Bayesian theorem

“Bayesian inference is the process of fitting a probability model to a set of data and summarizing the result by a **probability distribution on the parameters** of the model and on unobserved quantities such as predictions for new observations. “

Gelman, Bayesian Data Analysis.

$$\Pr(A|B) = \frac{\Pr(B|A) \Pr(A)}{\Pr(B)}$$

$$\Pr(disease|test\ positive) = \frac{\Pr(test\ positive|disease) \Pr(disease)}{\Pr(test\ positive)}$$

$$PPV = \frac{sensitivity * prevalence}{True\ positives + false\ positives}$$



Thomas Bayes (1763). An Essay Towards Solving a Problem in the Doctrine of Chances

Bayesian theorem & Bayesian inference

$$\Pr(\vartheta|y) = \frac{\Pr(y|\vartheta) \Pr(\vartheta)}{\Pr(y)}$$

Likelihood

= process model + observation model

→ likelihood of the observed data (y) under a model and parameter set (ϑ)

Prior

Suggestions for distributions for the parameter set ϑ BEFORE having seen the data y (“a-priori belief”)

Posterior

= Joint posterior probability distribution

→ Expresses uncertainty about the parameter set ϑ after taking both the data and the prior into account

Marginal likelihood (normalizing constant)


- Indicates what the data theoretically can look like BEFORE the actual data has been observed (using only the prior distribution and the log likelihood function).
- Calculated by integrating/summing the joint probability distribution over the parameter space
- Purpose: to normalize the posterior to a proper probability distribution (so that it integrates to 1)

“Prior”

- Encodes your “prior belief” of what the parameter distribution looks like
- Can and should influence the posterior
- The more data, the less influence has the prior
- Choose appropriate distribution:
- Beta() for proportions (0,1),
- Gamma() or LogNormal() for positive parameters
- Uniform() for uninformative priors without prior knowledge except the bounds

- Example:

- $\gamma \sim \text{Beta}(1.5, 10.0)$



Why Bayesian?

1) Bayesian is more intuitive:

- a) Bayesian credible interval: 95% probability, that the true estimate lies in this range
- b) Frequentist confidence interval: in 95% of cases, the true estimate lies within the specified range

2) Bayesian provides more information about the distribution of the parameters (rather than a point estimate and some CI

3) Bayesian acknowledges that there is uncertainty in parameters

4) Bayesian let's you update your model with new information (the prior)

The three steps of Bayesian inference

1) Setting up a full probability model:

- a) Define the process model
- b) Define an observation model
- c) Define the priors of the parameters of interest

2) Conditioning on observed data (aka fitting)

1) Evaluating the fit

- a) Evaluate the computation (convergence)
- b) Evaluate the fit to the data

Bayesian inference “by hand”: a practical example

Analytical solution vs. numerical solution

- Analytically calculating the marginal likelihood is difficult or even impossible for a large parameter set (requires integrals for continuous parameters)
- Calculating the profile log-likelihood for every possible value of the parameter is computationally very expensive
- We need a numerical (computational) approximation.
- We can use sampling strategies to sample (“draw”) from the posterior (rather than calculating it analytically)

Sampling algorithms: Markov Chain Monte Carlo MCMC

“Monte Carlo”

Class of methods for repeated random sampling from a distribution

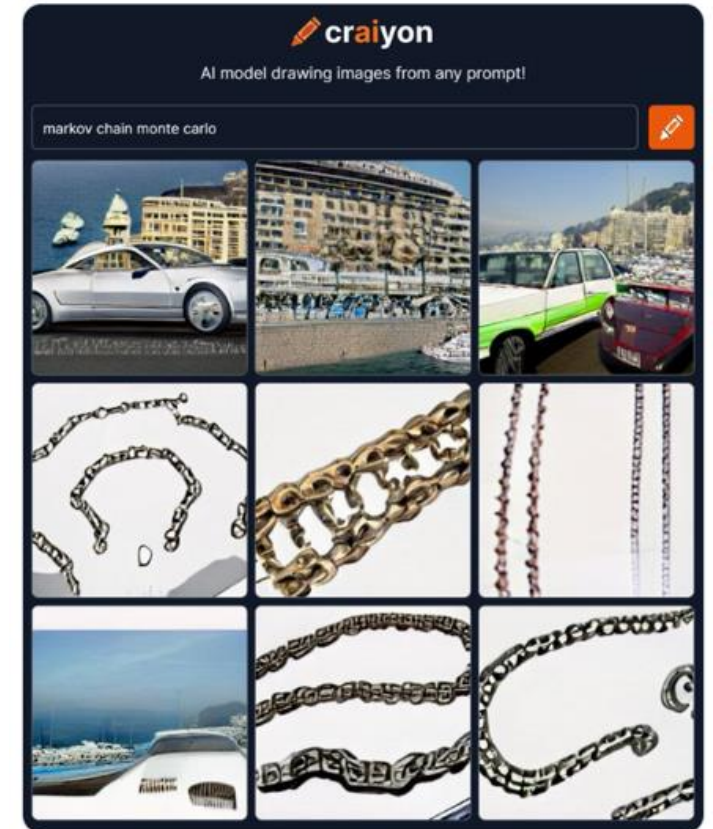
“Markov Chain”

A sequence of states, where each state depends only on the previous state

The “**MCMC**” algorithm generates a sequence (“markov chain”) of samples (“monte carlo”)

“Metropolis-Hastings MH”

A particular algorithm of using MCMC to sample from a distribution



Sampling algorithms: MH-MCMC


1. Start from an initial parameter sample θ and evaluate the posterior function (“current” target)
2. Define a proposal distribution: Multivariate Normal with mean= μ and pre-defined SD, draw a new parameter sample θ' from the proposal and evaluate the posterior function for the new θ' (“proposed target”)
3. Compare the current and proposed target function:

$$ratio = \frac{f(\theta')}{f(\theta)} = \frac{posterior(\theta')}{posterior(\theta)} = \frac{\frac{\Pr(y|\theta')\Pr(\theta')}{\Pr(y)}}{\frac{\Pr(y|\theta)\Pr(\theta)}{\Pr(y)}} = \frac{\Pr(y|\theta')\Pr(\theta')}{\Pr(y|\theta)\Pr(\theta)}$$

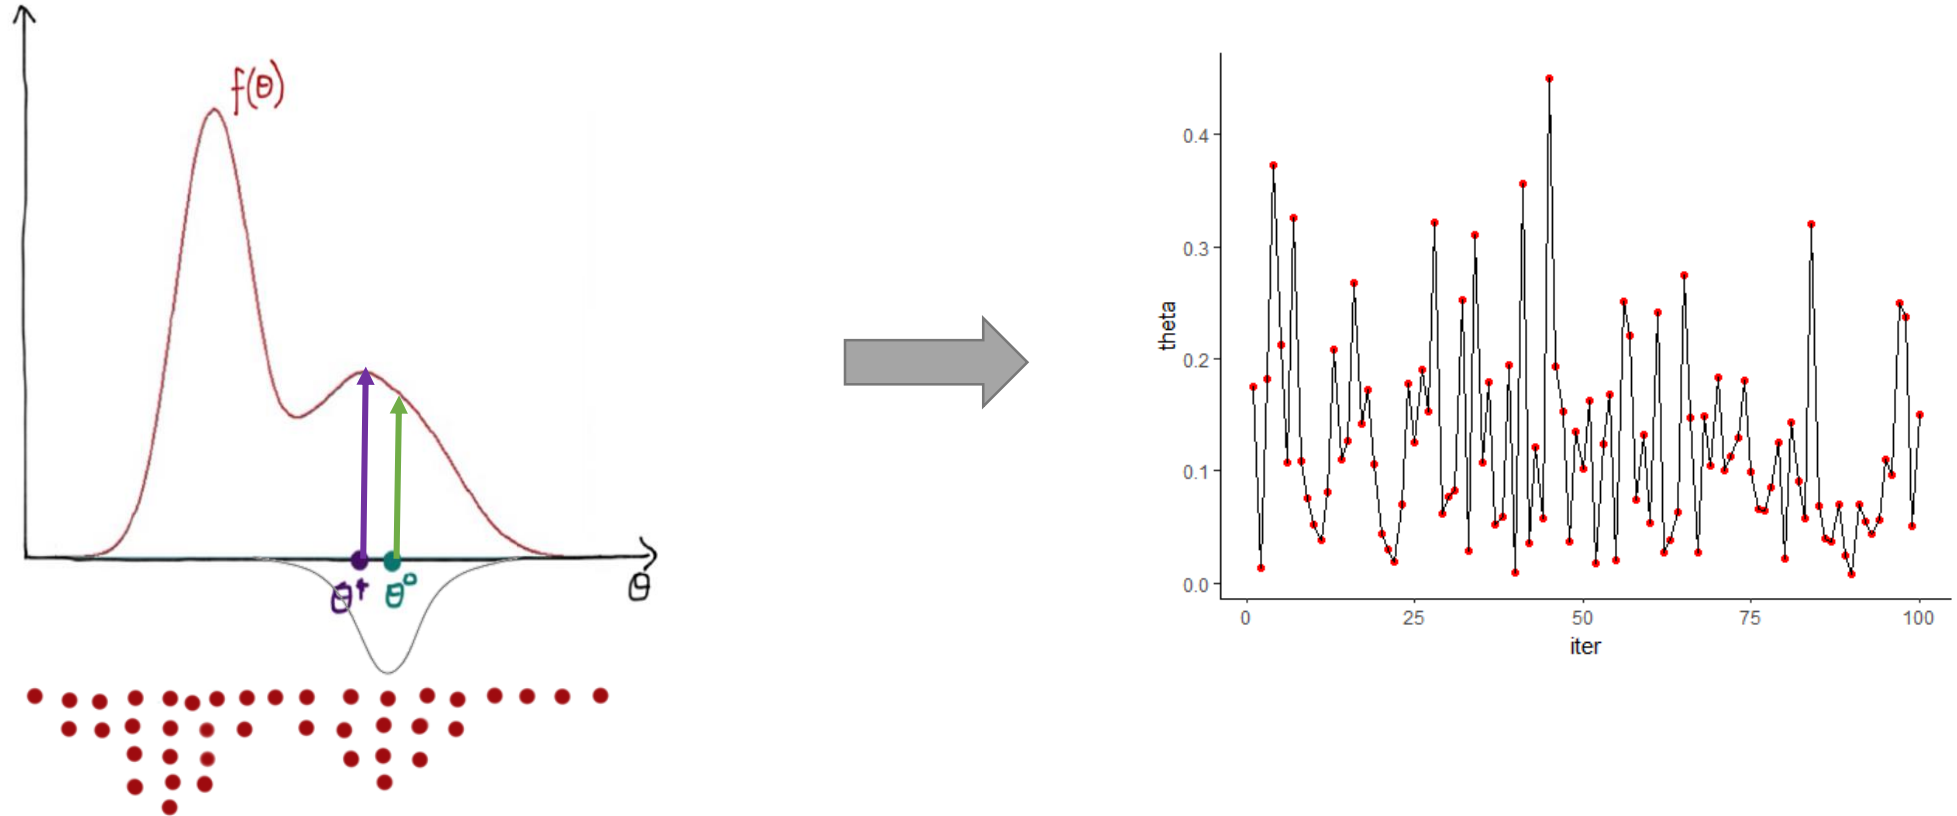
4. Calculate the acceptance ratio $r = \min(1, ratio)$
5. Draw a random uniform number $u(0,1)$
6. If $u < r \rightarrow$ accept θ' as the new current θ , else keep the current θ

Rinse and repeat 2-6 for n iterations

$$\frac{f(\theta')}{f(\theta)} = \frac{0.06}{0.05} = 1.2$$

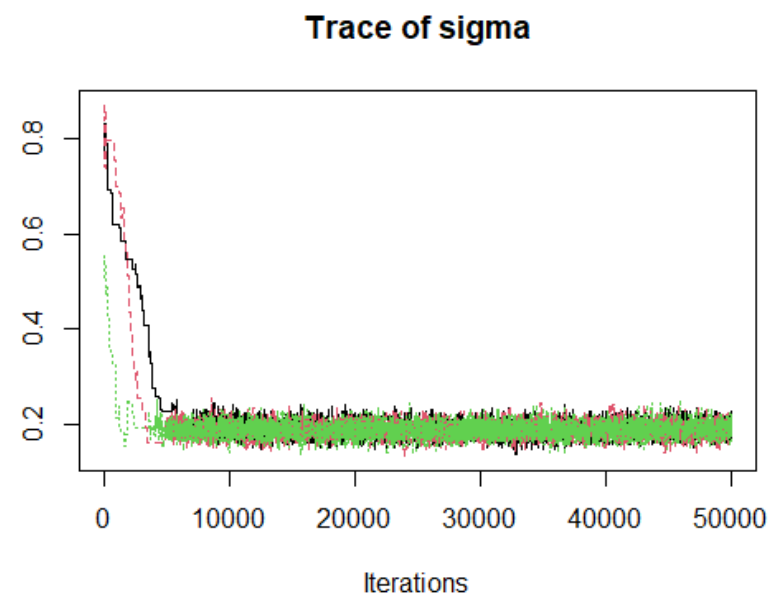





Sampling algorithms: MH-MCMC






With an increasing number of iterations (draws), the distribution of the proposed parameters (θ) more and more resembles the desired target (posterior) distribution
→ sampling from the posterior

Sampling algorithms: MH-MCMC



Convergence diagnostics


- - Convergence = Samples have reached a stationary distribution, which consists of samples of the posterior → fuzzy caterpillar
- Need multiple chains converging to the same posterior to fully confirm convergence
- Formal assessment of convergence with Rhat / PSRF: ratio of between and within chain variance. Rhat < 1.05 indicative of convergence
- Burn-in phase must be removed
- Aim for an acceptance rate of ~23% (for random walks)




Convergence diagnostics

Good chain mixing produces a fuzzy caterpillar


Proposal small \rightarrow small step size \rightarrow
high acceptance rate, poor mixing



Proposal large \rightarrow large step size \rightarrow
low acceptance rate, poor mixing



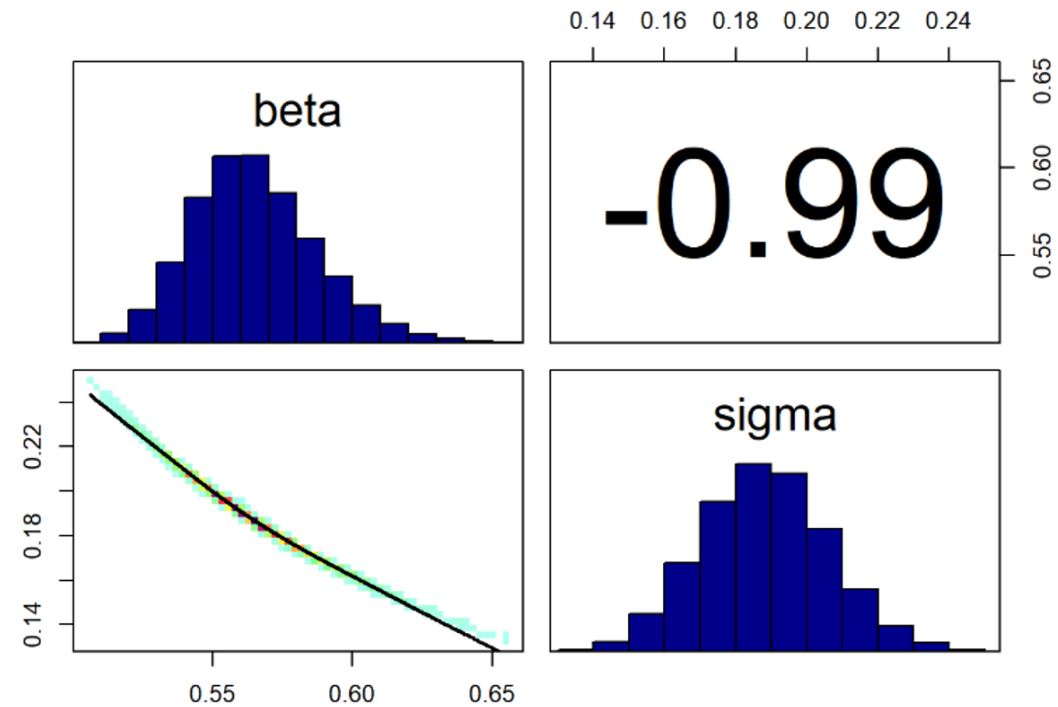

Optimal proposal \rightarrow hairy caterpillar



Effective sample size (ESS) indicates the number of independent samples (free from autocorrelation) \rightarrow should be at least 100.

Marginal posterior statistics



Visualize parameter correlation



Marginal posterior statistics

Summarize the posterior distributions with 95% credible intervals (contains the true parameter with 95% probability)


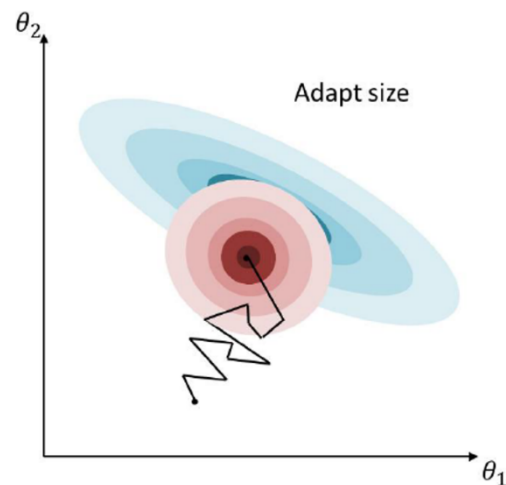
- 95% highest posterior density interval (HPDI) (narrowest interval that contains 95% of the data)
- Equal tailed 95% interval (2.5th and 97.5th quantile)



Adaptive MH-MCMC


- Standard MH-MCMC is very sensitive to choice of proposal (difficult to get good mixing)
- Poor performance with multiple parameters that are correlated
- Adaptive MH: alters the proposal during the burn-in to match the covariance matrix of the posterior. First step size (by varying the standard deviation), then shape (by varying the covariance between the parameter)

Two-stage adaptation



Sampling algorithms: DE-MCMCzs

Group of 3 chains, proposal for chain 1 is informed by the current and past state of the other two chains (no assumption of multivariate normal)



1. Choose current theta and evaluate the posterior

2. Project a vector to a randomly chosen state of another chain (current or past)

3. Select two other states randomly and project their distance onto the first projection line

4. The resulting difference vector is multiplied by some factor and added to the original point to draw the final proposal

5. The proposal is accepted with some probability

Rinse and repeat for all three chains

Choice of sampling algorithms

MH-MCMC


Adaptive MH-MCMC and other variants

Genetic algorithms (DE)


Parallel tempering

Hamiltonian Monte Carlo / NUTS

Increasing
model
complexity



The Bayesian workflow




Prior predictive checks

Prior predictive checks generate data from the prior distribution to check whether the prior is appropriate:

1. Sample multiple times from the prior
2. plug the values into the process model and simulate the trajectory
3. visualize the quantiles of the resulting trajectories together with the actual data

--> If the prior predictions do not overlap with actual data, you need to adjust the prior




Posterior predictive checks: trajectory fit

Posterior predictive checks aim to visualize the fit of the predicted trajectories to the data:

1. Sample multiple times from the posterior
2. plug the values into the process model and simulate the trajectory
3. visualize the quantiles of the resulting trajectories together with the actual data

--> If the posterior predictions do not overlap with actual data, you need to adjust the model, the prior or add/change the data



Trouble shooting

1. Choose a sampler appropriate for your model (HMC/NUTS/DE/parallel tempering >>> MH). Don't write your own sampler.
2. Follow the Bayesian workflow!
3. Start with the simplest model and the fewest parameters to fit. Gradually increase the complexity
4. Start by simulating data from the model and trying to fit the simulated model
5. Parameter correlation/non-identifiability makes convergence difficult:
 - Consider adding additional data (hospital data in addition to times series data) and a joint likelihood
 - Adjust/adapt/tighten your priors
 - Fix one of the correlated parameters if you can
6. Complex posteriors (local modes, abrupt changes) make convergence difficult:
 - Choose initial starting values for theta by sampling from the prior OR
 - Run MLE on the model first, and use the MLE as initial values for the MCMC

R packages for Bayesian ODE inference

Package	Samplers	Advantages	Disadvantages
LaPlacesDemon	Metropolis variants Genetic variants Gibbs HMC variants T-walk, SMC	<ul style="list-style-type: none">- Wealth of samplers to choose from- Fast!	<ul style="list-style-type: none">- Few example code published for ODE-IDD models
BayesianTools	Metropolis variants Genetic variants SMC T-walk	<ul style="list-style-type: none">- Requires to write out the log likelihood function and the prior functions (good for learning Bayes)	<ul style="list-style-type: none">- MH and DE not very performant for complex models (poor mixing, no convergence).- Not very fast
Stan/Rstan	HMC/NUTS	<ul style="list-style-type: none">- Very fast (C++)- Gradient-based samplers good for complex problems- Great documentation- Support from user forum	<ul style="list-style-type: none">- Stan DSL (awkward way of coding the model and the components)- Steep learning curve
pomp	pMCMC	<ul style="list-style-type: none">- Fast?	<ul style="list-style-type: none">- Pomp DSL (awkward way of coding the model)

Acknowledgments

- Sebastian Funk, LSHTM. Material from Bayesian Model fitting for infectious disease modelling → next course summer 2023
- Florian Hartig, Uni Regensburg. Material from Bayesian Model fitting with Bayesian Tools